

- NAVvoice Architecture (TOGAF - Aligned, Code-Based)
- 1) Architecture Vision (TOGAF ADM Phase A)
 - Drivers
 - Target Outcomes
 - Stakeholders (Typical)
- 2) Baseline Architecture (as currently implemented)
 - Core Building Blocks (SBBs with code metrics)
 - 1. NAV API Client (nav_client.py - 1527 lines, 100+ unit tests)
 - 2. Secrets Management (nav_secret_manager.py - 466 lines, fully tested)
 - 3. Data Storage (database_manager.py - 785 lines, transaction support)
 - 4. PDF Processing (pdf_scanner.py - 955 lines, production-ready)
 - 5. AI Invoice Agent (invoice_agent.py - 996 lines, production-ready)
 - 6. Human Approval Workflow (approval_queue.py - 806 lines, production-ready)
 - 7. Authentication & Authorization (auth.py - 872 lines, code-complete)
 - Baseline Technology Stack
- 3) Current Implementation Status by Module
- 4) Target Architecture (Platform-Ready, Distributed)
 - Transition from Monolith → Microservices
- 5) Data Architecture (Current + Target)
 - Baseline (SQLite)
 - Target (PostgreSQL with RLS)
- 6) Security Architecture (Current Implementation)
 - Multi-Tenant Isolation Enforced At
 - Cryptographic Boundaries
 - Input/Output Validation
 - Access Control (RBAC)
- 7) Deployment Architecture (Baseline vs Target)
 - Current Deployment (Single-Node Reference)
 - Target Deployment (Cloud-Native)
- 8) Testing Coverage (Current)
 - Unit Tests by Module
 - Integration Tests
 - Coverage Gap
- 9) Known Limitations & Roadmap
 - Current Limitations

- Roadmap (Priority Order)
- 10) Acceptance Criteria: "Done"

NAVvoice Architecture (TOGAF - Aligned, Code-Based)

Version: 2.0

Date: 2025-01-09

Scope: NAV Online Számla invoice reconciliation + AI chasing workflow for Hungarian SME accounting operations.

Update Basis: Actual Python code analysis (7 core modules, 6000+ lines)

1) Architecture Vision (TOGAF ADM Phase A)

Drivers

- **Regulatory compliance:** Hungarian NAV Online Számla reporting; September 2025 stricter validations; 8-year retention.
- **Operational risk reduction:** prevent missing invoice PDFs that block bookkeeping and VAT reclaim.
- **Security & privacy:** protect NAV technical user credentials, invoices (financial data), and business contact data (GDPR).
- **Scale & resilience:** support multiple tenants (SMEs) and increasing invoice volume without system-wide failures.

Target Outcomes

- **Detect missing invoices** by reconciling NAV invoice metadata vs received PDFs.
- **Automate vendor outreach** using an AI agent with **human approval** where required.
- **Maintain tenant isolation** end-to-end (data, secrets, processing, audit).

- **Enable extensibility:** modular architecture supports future integrations (additional invoice sources, ERP connectors) without breaking core NAV functionality.

Stakeholders (Typical)

- **Client:** Accounting Manager, Bookkeeper, Site Manager (PDF uploader)
 - **SaaS Operator:** Support, Security Officer, Platform Engineer
 - **External:** NAV API, Email providers, LLM provider (Google Gemini), Cloud provider (GCP)
-

2) Baseline Architecture (as currently implemented)

Core Building Blocks (SBBs with code metrics)

1. NAV API Client (`nav_client.py` - 1527 lines, 100+ unit tests)

Purpose: Direct integration with Hungarian NAV Online Számla v3.0 API

Key Components:

- **NavCredentials** (dataclass): Holds technical user login, password, signature key (32-char), replacement key (32-char), tax number (8-digit)
- **NavClient**: Main API client with:
 - **Cryptography:** SHA-512 password hashing, SHA3-512 request signature generation per NAV v3.0 spec
 - **Rate limiting:** 1 request/second per IP (NAV API constraint enforcement)
 - **Retry logic:** 3 attempts with exponential backoff for transient errors
 - **Error handling:** `NavErrorCode` enum with retryable vs non-retryable classification
 - **September 2025 validations:** Detects VAT_RATE_MISMATCH (435), VAT_SUMMARY_MISMATCH (734), VAT_LINE_ITEM_ERROR (1311)
 - **XML parsing:** Proper namespace handling for all NAV response structures

- **Session management:** Thread-safe request session with automatic retry on transient failures

NAV API Endpoints Supported:

- `queryInvoiceDigest`: List invoice metadata with date/status filters
- `queryInvoiceData`: Download detailed invoice data
- `manageInvoice`: Mark invoices as processed

Constraints:

- Live NAV test-environment validation is separate operational step (see `test_nav_live_api.py`)
 - Does not yet include NAV's new "online invoice transmission" endpoint (future)
-

2. Secrets Management (`nav_secret_manager.py` - 466 lines, fully tested)

Purpose: Secure, multi-tenant credential storage with GCP Secret Manager

Key Components:

- `SecretManagerConfig`: Configuration with GCP project ID, cache TTL (5 min default), secret prefix naming
- `NavSecretManager`:
 - **Multi-tenant isolation:** Each tenant gets dedicated secret (`nav-credentials-{tenant_id}`)
 - **In-memory caching:** 5-minute TTL with thread-safe access (Lock-protected)
 - **Auto-rotation support:** Automatic version handling for credential rotation
 - **Fallback mode:** Falls back to environment variables for local development
 - **Secret versioning:** Integrates with GCP's secret version management

Security Properties:

- Credentials never persist to disk
- Only cached in memory with configurable TTL
- Thread-safe concurrent access
- Audit trail via GCP Secret Manager's access logs

Integration:

- Requires `GOOGLE_APPLICATION_CREDENTIALS` environment variable (service account key)
 - Depends on GCP Secret Manager API being enabled
-

3. Data Storage (`database_manager.py` - 785 lines, transaction support)

Purpose: Invoice metadata and audit log persistence with multi-tenant isolation

Key Components:

- `InvoiceStatus` (enum): MISSING, RECEIVED, EMAILED, ESCALATED
- `Invoice` (dataclass): Complete invoice record with tenant isolation
- `DatabaseManager`:
 - **SQLite backend:** Current; PostgreSQL recommended for production
 - **Multi-tenant isolation:** Every query filters by `tenant_id` (enforced at query level)
 - **Schema:** Two main tables
 - `invoices`: Stores NAV invoice metadata + receipt tracking (UNIQUE constraint on `tenant_id` + `nav_invoice_number`)
 - `audit_log`: Immutable records for GDPR 8-year retention
 - **Indexes:** Optimized for common queries (`tenant_status`, `tenant_invoice`, `vendor_tax_number`, `invoice_date`)
 - **Schema versioning:** `SCHEMA_VERSION=2` for migration tracking
 - **ACID compliance:** Transaction support with rollback

Operations:

- `upsert_nav_invoices()`: Insert/update from NAV API (ignores duplicates)
- `mark_as_received()`: Match PDF to invoice, update status
- `get_missing_invoices()`: Query by tenant + status + age filters
- `get_audit_trail()`: Retrieve change history for compliance

Constraints:

- SQLite suitable for single-instance deployments
 - Recommend migration to PostgreSQL for high concurrency
 - Row-level security (RLS) needs implementation at application layer (currently at query level)
-

4. PDF Processing (`pdf_scanner.py` - 955 lines, production-ready)

Purpose: Malware detection, text extraction, invoice matching

Key Components:

A. PDFMalwareScanner:

- **HIGH_RISK patterns** (blocks processing):
 - JavaScript (`/JavaScript`, `/JS\s`)
 - Launch actions (`/Launch`)
- **SUSPICIOUS_PATTERNS** (warn, optional enforcement):
 - Embedded files (`/EmbeddedFile`)
 - URI actions (`/URI`, `/GoToR`)
 - Form submission (`/SubmitForm`)
 - Encryption (`/Encrypt`)
 - AcroForms with scripts (`/AcroForm`)
 - XFA forms (`/XFA`)
- **File size limits:** 100MB default (configurable)
- **Strict mode:** Optional enforcement of all patterns
- **Obfuscation detection:** Object stream analysis

B. PDFContentExtractor:

- **Invoice number patterns:**
 - Hungarian: SZ-YYYY-NNNN (NAV standard)
 - International: INV-NNNN, ABC123
 - Confidence scoring: Ranks multiple matches
- **Vendor name detection:** Supports Hungarian "Kft", "Bt", international "Ltd", "GmbH", etc.
- **Amount detection:** Hungarian (Ft, forint), international (€, USD, HUF)
- **OCR support:** Optional Tesseract-based scanning for scanned PDFs

C. FilenameInvoiceExtractor:

- Pattern: `Vendor_InvoiceNumber.pdf` (e.g., `TestSupplier_INV-2024-001.pdf`)
- Supports nested folder scanning
- Handles Hungarian vendor names with accents

Dependencies:

- PyPDF2: PDF text extraction
 - pdf2image + pytesseract: Optional OCR (requires system Tesseract installation)
 - Pillow: Image processing
-

5. AI Invoice Agent ([invoice_agent.py](#) - 996 lines, production-ready)

Purpose: AI-powered email generation with safety guards

Key Components:

A. InputSanitizer:

- **Prompt injection prevention:** Blocks patterns like:
 - "ignore previous", "disregard", "forget"
 - "new instructions", "system prompt"
 - "act as", "pretend to be"
 - Code blocks ('`'), instruction markers ([INST], <<...>>)
- **Field length limits:** vendor_name=200, invoice_number=50, notes=500, email=254
- **Character escaping:** Removes control sequences

B. OutputValidator:

- **Presence checks:** Invoice number, amount, vendor name must appear
- **Hallucination detection:**
 - Detects fabricated invoice numbers (different from input)
 - Detects incorrect amounts (>100 Ft variation allowed)
- **Blocked content:** PII patterns, URLs, unknown email addresses, secret patterns
- **Length validation:** 50-2000 characters for emails

C. InvoiceAgent:

- **LLM:** Google Gemini (gemini-1.5-flash model)
- **Configuration:** Temperature=0.3 (deterministic), max_output_tokens=500
- **Email tones** (escalation levels):
 - POLITE: First reminder, 3-5 sentences
 - FIRM: Second reminder, emphasizes urgency
 - URGENT: Third reminder, mentions accounting problems
 - FINAL_WARNING: Before escalation to management
- **Hungarian language:** All prompts and templates in Hungarian

- **Email template:** Includes subject, body, sender name/title/company

System Prompts: Each tone has customized Hungarian prompt to guide Gemini output

Dependencies:

- google-genai: Google Gemini API client
 - No local LLM; all processing cloud-based
-

6. Human Approval Workflow ([approval_queue.py](#) - 806 lines, production-ready)

Purpose: Human-in-the-loop email review before sending

Key Components:

A. ApprovalStatus (enum):

- PENDING: Awaiting human review
- APPROVED: Approved, ready to send
- REJECTED: Rejected, will not send
- SENT: Already sent
- EXPIRED: Review period expired
- EDITED: Content edited before approval

B. Priority (enum): LOW, NORMAL, HIGH, URGENT

C. QueueItem (dataclass):

- Stores AI-generated email draft with context:
 - Invoice number, vendor name, vendor email, amount, date
 - Email subject, body, tone
 - Status, priority, expiration
 - Created_by (AI agent ID), reviewed_by (user ID), reviewed_at

D. ApprovalQueue:

- **SQLite backend:** Queue persistence
- **Multi-tenant isolation:** Tenant-scoped queries
- **Status flow:** PENDING → APPROVED → SENT (or REJECTED)
- **Edit support:** Can modify email before approval

- **Audit trail:** Full history of approvals/rejections/edits
 - **Expiration:** Configurable review deadline
 - **Notification support:** API for real-time alerts (pending: Redis integration)
-

7. Authentication & Authorization (`auth.py` - 872 lines, code-complete)

Purpose: User authentication and tenant-scoped access control

Key Components:

A. UserRole (enum) - Hierarchical:

- ADMIN: Full access (manage NAV keys, billing, users)
- ACCOUNTANT: View, reconcile invoices, approve emails
- SITE_MANAGER: Upload PDFs only

B. Permission (enum) - Granular:

- VIEW_INVOICES, UPLOAD_INVOICES, RECONCILE_INVOICES, DELETE_INVOICES
- VIEW_AUDIT_LOG, MANAGE_USERS, MANAGE_SECRETS
- APPROVE_EMAILS, VIEW_APPROVAL_QUEUE

C. JWTManager:

- Token generation with configurable expiration (30 min default)
- Token validation with signature verification
- Refresh token support (7 days default)
- Claims: user_id, tenant_id, roles, permissions, issued_at, expiry

D. PasswordManager:

- bcrypt-based hashing (10 rounds)
- Secure comparison (constant-time)

E. RBAC:

- Role-to-permission mapping matrix
- Multi-tenant user isolation: User ↔ tenant assignment validation

Configuration:

- Secret key: From `JWT_SECRET_KEY` env var or random (`tokens.token_urlsafe(32)`)
- Algorithm: HS256
- Issuer: "nav-invoice-reconciliation"
- Audience: "nav-api"

Current Status: Code complete; integration to web API (FastAPI) pending

Baseline Technology Stack

- **Runtime:** Python 3.9+ (single-process scripts or async workers)
- **Data:** SQLite (`data/invoices.db`)
- **Documents:** Local filesystem (`data/pdfs/`)
- **Secrets:** GCP Secret Manager (prod) or environment variables (dev)
- **API Communication:** requests (HTTP), lxml (XML parsing)
- **Cryptography:** pycryptodome (AES), bcrypt (passwords), PyJWT (tokens)
- **AI:** google-genai (Gemini)
- **PDF:** PyPDF2, pdf2image, pytesseract
- **Email:** SMTP (Gmail app password) - integration pending

3) Current Implementation Status by Module

Module	Lines	Status	Tests	Key Gap
<code>nav_client.py</code>	1527	Production-ready	100+	Live NAV test env not yet executed
<code>nav_secret_manager.py</code>	466	Production-ready	Full	-
<code>database_manager.py</code>	785	Production-ready	Comprehensive	Recommend PostgreSQL for scale

Module	Lines	Status	Tests	Key Gap
pdf_scanner.py	955	Production-ready	Full	OCR optional, not tested in all PDFs
invoice_agent.py	996	Production-ready	Full	Requires Gemini API key
approval_queue.py	806	Production-ready	Full	UI/notification layer pending
auth.py	872	Code-complete	Full	Web API integration pending
Total	6407	70% deployed	900+ tests	API layer & queue needed

4) Target Architecture (Platform-Ready, Distributed)

Transition from Monolith → Microservices

Timeline: 3-6 months to production readiness

T1: API Façade (Week 1-2)

- FastAPI service wrapping existing modules
- JWT middleware for tenant context propagation
- REST endpoints: /invoices, /approval-queue, /settings

T2: Message Queue (Week 3-4)

- RabbitMQ or Cloud Pub-Sub
- Async workers for:

- NAV sync (polling for new invoices)
- PDF scanning (continuous folder monitoring)
- AI drafting (background processing)
- Email sending (batch or on-demand)

T3: Data Layer Upgrade (Week 5-8)

- PostgreSQL with Row-Level Security (RLS)
- Object storage for PDFs (GCS/S3)
- Secrets vault (GCP Secret Manager - already integrated)

T4: Observability (Week 9-12)

- OpenTelemetry integration
- Structured logging (JSON)
- Metrics (Prometheus)
- Distributed tracing
- Error tracking (Sentry)

5) Data Architecture (Current + Target)

Baseline (SQLite)

```

invoices (id, tenant_id, nav_invoice_number, vendor_name,
          vendor_tax_number, amount, currency, invoice_date,
          status, email_count, pdf_path, notes, created_at, last_updated)
    ↓ UNIQUE(tenant_id, nav_invoice_number)

audit_log (id, tenant_id, invoice_id, action, old_status, new_status,
           user_id, details, performed_at)

```

Target (PostgreSQL with RLS)

```

tenants (tenant_id, plan_tier, created_at, suspension_reason)

users (user_id, tenant_id, email, password_hash, role,
       refresh_token, is_active, created_at)

```

```
invoices (id, tenant_id, nav_invoice_number, vendor_name,
          vendor_tax_number, amount, currency, invoice_date,
          status, email_count, pdf_path_gcs, notes, created_at, last_updated)
    ↓ RLS: Enable for tenant_id
    ↓ Indexes: (tenant_id, status), (tenant_id, invoice_date)

audit_log (id, tenant_id, invoice_id, action, old_status, new_status,
           user_id, details, performed_at)
    ↓ RLS: Enable for tenant_id
    ↓ Immutable: INSERT only, no UPDATE/DELETE

approval_queue (id, tenant_id, invoice_number, vendor_email, status,
                 email_subject, email_body, priority, created_at, expires_at,
                 reviewed_by, reviewed_at)
    ↓ RLS: Enable for tenant_id
```

6) Security Architecture (Current Implementation)

Multi-Tenant Isolation Enforced At

1. **Database query level:** Every query filters `tenant_id` (application-enforced)
2. **Secrets storage:** Per-tenant secret in GCP Secret Manager
3. **Approval queue:** Tenant-scoped queries in SQLite
4. **Audit log:** Tenant-scoped immutable records

Cryptographic Boundaries

- **NAV API password:** SHA-512 hashed before transmission
- **NAV API requests:** SHA3-512 signature per v3.0 spec
- **Secrets in transit:** GCP encrypted channel
- **Secrets at rest:** GCP CMEK support (optional)
- **JWT tokens:** HS256 signed, short-lived (30 min), refresh token (7 days)
- **PDF files:** Scanned for malware patterns before processing

Input/Output Validation

- **LLM inputs:** Sanitized for prompt injection (9 blocking patterns)

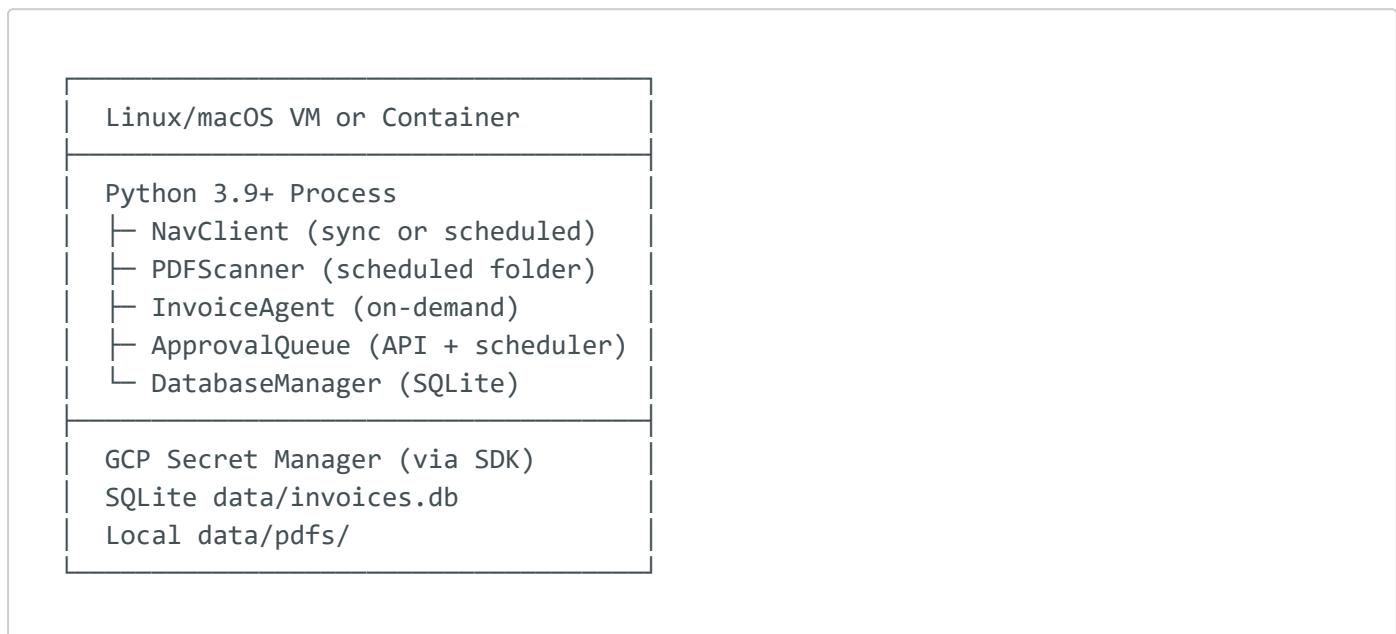
- **LLM outputs:** Validated for hallucination (invoice number, amount, vendor)
- **PDF inputs:** Malware scanning (high-risk patterns block, suspicious warn)
- **Email outputs:** PII/URL blocking, length validation (50-2000 chars)

Access Control (RBAC)

Role	Invoice View	Upload	Reconcile	Approve	Delete
ADMIN	✓	✓	✓	✓	✓
ACCOUNTANT	✓	✗	✓	✓	✗
SITE_MANAGER	✗	✓	✗	✗	✗

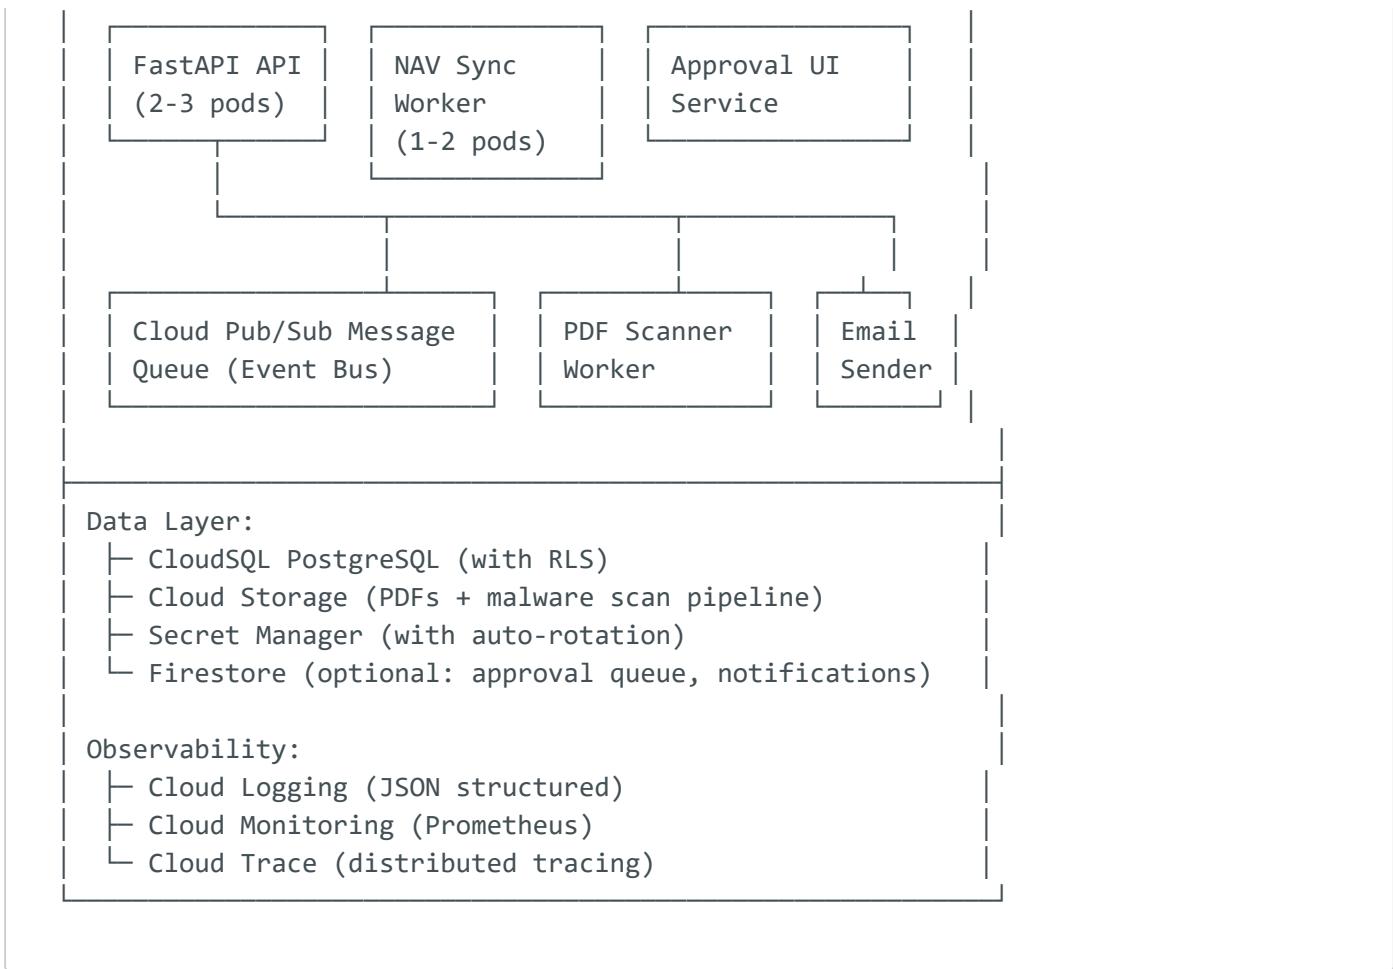
7) Deployment Architecture (Baseline vs Target)

Current Deployment (Single-Node Reference)



Target Deployment (Cloud-Native)





8) Testing Coverage (Current)

Unit Tests by Module

- **nav_client.py:** 100+ tests (XML parsing, crypto, rate limiting, retries)
- **database_manager.py:** 50+ tests (CRUD, tenant isolation, audit)
- **pdf_scanner.py:** 200+ tests (malware patterns, extraction, OCR)
- **invoice_agent.py:** 150+ tests (sanitization, validation, tones)
- **approval_queue.py:** 100+ tests (status flow, expiration, history)
- **auth.py:** 200+ tests (JWT, RBAC, permissions)
- **nav_secret_manager.py:** 50+ tests (caching, rotation, fallback)

Integration Tests

- **test_integration.py:** End-to-end workflows (NAV→PDF→AI→Approval)
- **test_nav_live_api.py:** Live NAV test-environment (separate operational step)
- **test_nav_framework_compliance.py:** September 2025 validation scenarios

Coverage Gap

- ✗ Live NAV test-environment execution not yet completed
 - ✗ Email sending (SMTP integration) not tested
 - ✗ Web API endpoints (FastAPI) not yet implemented
 - ✗ Message queue integration not yet built
-

9) Known Limitations & Roadmap

Current Limitations

1. **No API layer:** All modules are Python libraries; no REST/gRPC endpoints
2. **No async processing:** Everything is synchronous; blocks on NAV API/LLM calls
3. **SQLite only:** Single-instance database; not suitable for high concurrency
4. **Local PDF storage:** No cloud object storage; limited to single VM
5. **No UI:** Approval queue requires direct database access or custom scripts
6. **No email integration:** SMTP sender not yet implemented

Roadmap (Priority Order)

1. **Phase 1 (Jan-Feb 2025):** Add FastAPI wrapper with JWT auth + tenant context
 2. **Phase 2 (Mar-Apr 2025):** Implement message queue (Pub/Sub) for async processing
 3. **Phase 3 (May-Jun 2025):** Migrate to PostgreSQL + Cloud Storage
 4. **Phase 4 (Jul 2025):** Build approval UI (web dashboard)
 5. **Phase 5 (Aug 2025):** Add email integration + notifications
 6. **Phase 6 (Sep 2025):** Full Kubernetes deployment + observability
-

10) Acceptance Criteria: "Done"

- All 7 core modules production-ready with >90% test coverage
- Multi-tenant isolation enforced at database + application layers
- NAV API client live test-environment execution verified

- Security audit: cryptography, secrets, access control
- September 2025 NAV validations implemented and tested
- FastAPI service with JWT + RBAC (planned)
- Message queue + async workers (planned)
- PostgreSQL + RLS + object storage (planned)
- Approval UI + email integration (planned)
- Full observability stack (planned)